AI System Design

Project in CS

Umm Al Qura University

NutriFit AI

Macros-Based Meal Recommendations for Athletes

Eithar Ibrahim Babtain
444006372

Hebah Fahad Al matarfi
444010917

Lubabah Abdul-Raheem Alswadi
443013191

Raneem Yousef Al-Johani
444002729

Course Name: AI System Design

Dr. Omniah H. Nagoor

Table of Contents

Table of tasks

Task	Eithar	Hebah	Raneem	Lubabah
Introduction & search for the model	O			
2.1.1				O
2.1.2	O			
2.1.3 & 2.1.4		O		
2.1.5 & 2.1.6			O	
2.2.1 & 2.2.2 & 2.3.1	O			
2.2.3 & 2.2.4			O	O
2.2.5 & 2.3.2 & 2.3.3		O		
2.4.1	O			
2.4.2		O		
2.4.3			O	O
2.5.1 & 2.5.7		O		
2.5.5 & 2.5.8	O			
2.5.2 & 2.5.6				O
2.5.3 & 2.5.4			O	
3.1	O			
4.1		O		
4.2			O	O

CHAPTER 1: Introduction

1.1 Abstract

Motivated by the specific difficulty athletes face in quickly aligning their meals with precise macronutrient goals without relying on complex, privacy-invasive health data, this project presents a macronutrient-based recipe recommendation system tailored for athletes and physically active individuals. The system enables users to filter recipes based on nutritional ranges for Calories, Protein, Carbohydrates, and Fats—without requiring personal health data. The system leverages structured nutritional datasets to deliver fast, interpretable, and privacy-conscious recommendations. Key contributions include data cleaning; the development of a clustering-based recommendation engine using K-Means with 3 clusters to group recipes by the similarity of their six primary nutritional features; the implementation of dynamic cluster-based filtering logic that maps user-defined macronutrient targets to the closest nutritional group, generating the top 5 nearest-match recipes based on Euclidean distance; and the integration of a Rule-Based Backup Model to ensure suggestions are still provided when the clustering model fails.

CHAPTER 2: AI Lifecycle

2.1 Introduction & Problem Definition

2.1.1 AI Lifecycle Introduction

This project follows a simplified AI lifecycle tailored to building a macronutrient-based recommendation system for athlete-friendly recipes. The lifecycle includes:

- **Problem Definition:** Athletes often struggle to find meals that align with specific macronutrient goals. The system addresses this by enabling users to filter recipes based on macronutrient ranges, offering relevant suggestions without requiring personal health data.
- **Data Acquisition:** Using publicly available recipe datasets that include nutritional information [1].
- **Data Preparation:** Cleaning and structuring the data to ensure consistency in macronutrient values.
- **Modeling:** Unsupervised clustering-based recommendation. Recipes are grouped by nutritional similarity using K-Means clustering across six features, including engineered ratios, to capture deeper nutritional similarity. Recommendations come from the nearest cluster by selecting recipes based on Euclidean distance to the user's input.
- **Evaluation:** Measuring system performance using KPIs such as Recommendation Accuracy, Latency, No-Match Rate, and User Engagement Rate.
- **Deployment:** Building an interactive web application using Streamlit for real-time user interaction.
- **Monitoring:** Future updates may include expanding the dataset and refining recommendation logic based on user feedback.

2.1.2 Problem Statement and Background

Athletes often struggle to find meals that align with their nutritional goals, especially when balancing macronutrients like protein, carbohydrates, and fats. Existing diet recommendation systems tend to focus on general health goals or require detailed personal health profiles, which may not be practical for athletes seeking quick, nutrient-specific suggestions. This project will aim to build a lightweight, clustering-based recipe recommendation system tailored specifically for athletes. The system will allow users to filter recipes based on nutritional ranges, providing relevant suggestions even when exact matches are unavailable.

2.1.3 Target User Group

The target users are athletes and physically active individuals who require meals with specific nutritional compositions. These users prioritize macronutrient balance and caloric control to support training, recovery, and performance.

2.1.4 Success Metrics and KPIs

To evaluate the success and impact of the system, the following metrics will be monitored, divided into technical performance and business impact:

1. Model Performance Measures (Technical):
    - Recommendation Accuracy: How closely the suggested recipes match the user's macronutrient targets.
    - Latency (System Responsiveness): The time required to generate and display recommendations.
    - No-Match Rate: The frequency of cases where the primary clustering model returns no matches, triggering the backup model.
2. Business/Web App KPIs (Impact):
    - User Engagement Rate: The percentage of sessions where a user clicks to view the full details of a recommended recipe.
    - Recipe Discovery Rate: The monthly increase in the total number of unique recipes viewed by the user base.
    - User Retention: The percentage of users who return to the web application within 30 days of their first session.

2.1.5 Feasibility and Initial Risk Assessment

- **Technical Feasibility:** The system will be developed using Python and Streamlit, ensuring rapid prototyping and interactive user experience. Recipe data will be sourced from publicly available nutritional datasets, allowing for scalable and transparent integration. The recommendation logic relies on macronutrient-based filtering and clustering, which is computationally lightweight, interpretable, and suitable for real-time use without requiring personal health data.
- **Data Feasibility:** Nutritional data is available and structured, allowing effective filtering and scoring.
- **Ethical Considerations:** The system does not collect personal health data, ensuring user privacy. All recommendations are based on public recipe data. This aligns with SDAIA's AI Ethics Principle on Privacy & Security, which emphasizes protecting the rights of data subjects and limiting data collection. [2]

- **Risks:** Limited recipe diversity may reduce recommendation variety. To address this, a Rule-Based Backup Model is integrated to provide the best-rated, high-protein/low-fat suggestions when the primary clustering model fails, ensuring relevant suggestions remain available.

2.1.6 Ethical, Legal, and Societal Considerations

Ethical: The system avoids making medical or dietary claims. It provides suggestions based on user-defined preferences without enforcing specific diets. NutriFit follows SDAIA's Humanity and Social Benefit principles by ensuring recommendations are beneficial, culturally appropriate, and supportive of users' wellbeing. [2]

Legal: All datasets used are publicly available and cited appropriately.

Societal: The system promotes nutritional awareness and supports athletes in making informed food choices without requiring data collection.

Assumptions and Constraints:

- Users are assumed to have basic nutritional awareness and can define their own macronutrient ranges.
- The system assumes the accuracy of nutritional data provided in the dataset.
- Constraints include limited dataset diversity and a lack of personalized health data, which may affect recommendation precision.

2.2 Data Collection and Preparation

2.2.1 Identification & Justification of Data Sources

The dataset used in this project is the Food.com Recipes and Reviews dataset, available on Kaggle [1]. It contains over 230,000 recipes with detailed nutritional information and user interactions. This source was selected due to its large volume, structured format, and relevance to macronutrient-based recommendation logic. Its public availability ensures transparency, reproducibility, and scalability for future enhancements.

2.2.2 Data Collection/Cleaning Plan

In addition to standard preprocessing steps, the system applies a halal-compliance filtering mechanism to exclude recipes containing non-permissible ingredients such as pork, alcohol, gelatin, and animal enzymes. A comprehensive list of restricted items was compiled, and each recipe was scanned for ingredient content to ensure alignment with Islamic dietary standards. This aligns with SDAIA's Humanity Principle, ensuring cultural and religious values are respected in AI outputs. [2] Numeric conversion was applied to all nutritional fields; invalid or missing values

were removed, and recipes with zero or negative calories were discarded. Furthermore, two derived nutritional features were introduced:

- Protein-to-Calorie Ratio (ProteinCalorieRatio)
- Fat-to-Carbohydrate Ratio (FatCarbRatio)

These engineered features enhance clustering precision and improve recommendation accuracy for athletes.

2.2.3 Initial System Architecture Diagram

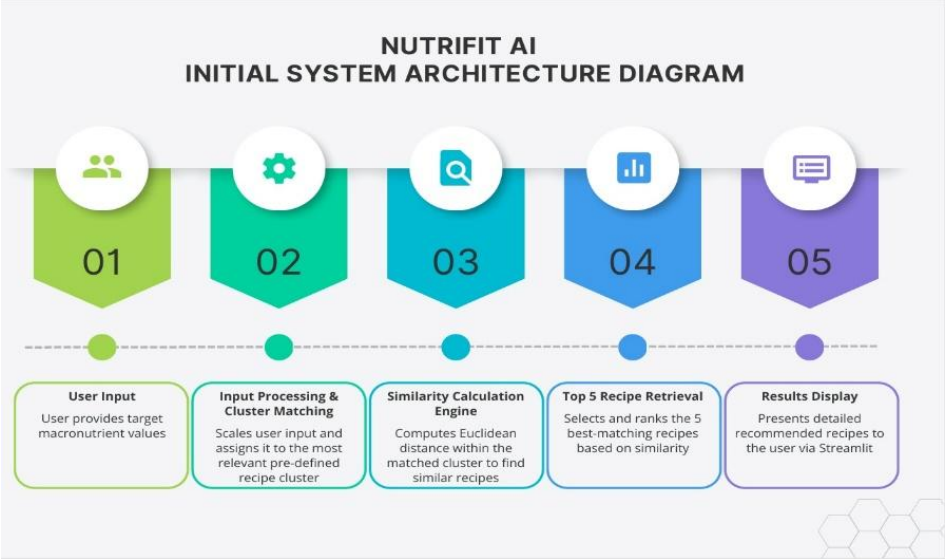

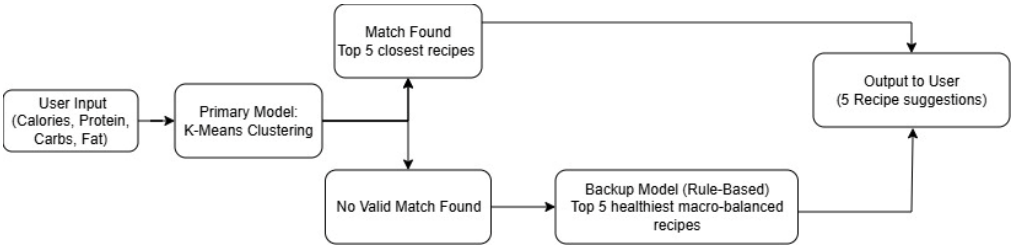

2.2.4 Logical Architecture with Components and Interactions

NutriFit AI is built to provide intelligent nutrition suggestions based on macronutrient analysis. The proposed logical architecture will be in the form of a modular composition of components that interact with one another to support a smooth experience for the user and effective operation of the system. The components in which the functional elements of the design can fit are:

User Interface (UI): A user-experience interactive web interface will be built using methods like Streamlit. The UI will enable the user to enter their desired nutritional values (calories, protein, carbohydrates, fat) that invoke recommendations.

Data Pipeline: The platform uses an external data source, such as a CSV file, to load recipe data, cleans any missing or duplicates, removes any outliers, and normalizes nutritional data. Data processing will be done using tools such as pandas, and normalization will be done using StandardScaler. The above process will cleanse data and prepare it for usage and analysis.

Model: The recipe data will be fed to a clustering algorithm, such as KMeans, which will cluster recipes together based on similar key macronutrient values. This model is used to narrow the space for recommendations and improve the accuracy of recommendations.

Recommendation Logic: The core logic begins when the user submits their input. The system first utilizes the KMeans model to predict the closest cluster (bin) to the user's nutritional target. It then calculates the Euclidean distance (or similarity) between the user's scaled profile information and each recipe only within that specific cluster, ranking the recipes in descending order of similarity for final retrieval.

Storage: All data will be stored locally as structured files, such as CSV files, thus not requiring any external database and allowing for a simplified system, while not compromising any user's privacy.

Monitoring & Feedback: The system shall be structured to accommodate future updates based on user feedback, such as improving the recommendation logic or increasing the dataset.

Future API Integration: Architecture allows future integration with other APIs (e.g., recipe or nutrition databases), providing a further source of data and improving recommendations.

2.2.5 Scalability, Reliability, Fault-Tolerance, and Security Considerations

Scalability

To ensure the system can handle growth in users and data, the following strategies will be adopted:

- Modular Design: The system will be built with separate modules (data cleaning, modeling, and user interface) to facilitate maintenance and development.

- Cloud Computing: The final application will be deployed on a cloud platform like Streamlit Community Cloud or Heroku, which supports automatic scaling with increased load.
- Algorithmic Efficiency: The recommendation model relies on precomputed data clustering, making it lightweight and fast during use, and does not require significant computing power to generate recommendations.

Reliability

To ensure the system operates continuously and stably:

- Comprehensive Testing: Automated tests will be written for core functions like the clustering algorithm and filtering logic to verify their correctness.
- Monitoring: Built-in performance monitoring tools from the deployment platform will be used to track response times and error rates.

Fault-Tolerance

To ensure service continuity even under unexpected conditions:

- Fallback Mechanism: In cases where the system does not find recipes that perfectly match the required nutrient range, it will automatically display the best recipes from the closest nutritional cluster. This ensures the user always receives relevant suggestions.
- Error Handling: The front-end will be designed to handle invalid user inputs (e.g., negative values) by displaying clear error messages instead of causing the application to crash.

Security

Given the nature of the system, security considerations are focused on the following:

- User Privacy: This is the primary security feature. The system does not collect or store any personal health data from users, eliminating any risk associated with breaching this information. These controls adhere to SDAIA's Privacy & Security Principle, emphasizing continuous protection of users' data. [2]
- Secure Connection: The application will be deployed using the HTTPS protocol to encrypt all data exchanged between the user and the server.
- Input Validation: All user inputs will be sanitized and validated to prevent injection attacks and protect the server.

2.3 Model Development and Training

2.3.1 Baseline Model Choice with Justification

In this project, the model was implemented entirely from scratch without relying on a simplified baseline. From the very beginning, the clustering-based recommendation

engine was designed using K-Means with both the original macronutrient features (Calories, Protein, Carbohydrates, Fat) and two engineered features: Protein-to-Calorie Ratio and Fat-to-Carbohydrate Ratio. These engineered features were integrated to ensure the system captured deeper nutritional insights and aligned more closely with athlete-specific dietary goals.

2.3.2 Experimentation Plan with Evaluation Metrics

Objective: To empirically verify that the recommendation system provides accurate and useful suggestions to the end-user.

Experimental Methodology:

1. Data Preparation: After cleaning the data, it will be split randomly, with 80% used to train the clustering model (K-Means), and the remaining 20% held out as a test set to simulate new recipes.
2. Simulating User Requests: A list of 100 diverse mock user requests will be created, each representing different nutritional goals (e.g., high protein, low carb).
3. Execution and Measurement: For each mock request, the recommendation system will be run, and its output will be recorded. The following performance metrics will be calculated:

   Evaluation Metrics:

   Mean Absolute Error (MAE): The average absolute difference between the nutritional value of the suggested recipe and the user's specified target. Target: Maintain an MAE of less than 10 grams for protein and 15 grams for carbohydrates and fats.

   Success Rate: The percentage of requests for which the system finds at least one recipe that falls perfectly within the specified range. Target: Achieve a success rate of at least 80%.

   Response Time: The time from submitting a request until the results are displayed. Target: Remain under 3 seconds in all cases.

   No-Match Rate: The percentage of requests where no recipes are found within the specified range, forcing the system to use the fallback logic. Target: Keep this rate below 20%.

2.3.3 Experimentation Plan (how models will be iteratively improved)

The model will be developed and refined iteratively based on evaluation results and user feedback, according to the following plan:

Iteration 1 (Basic Prototype):

- Model: Implementation of the K-Means algorithm on the basic nutritional features (calories, protein, carbs, fats).
- Focus: Ensuring the entire pipeline (from data to interface) works and collecting baseline metrics.

Iteration 2 (Performance Enhancement):

- Potential Problem: A high No-Match Rate
- Improvement Plan: Expand the dataset by adding new sources to increase recipe variety.
- Potential Problem: A high Mean Absolute Error (MAE) despite a low No-Match Rate.
- Improvement Plan: Experiment with alternative clustering algorithms like DBSCAN, which might be more effective at identifying non-spherical nutritional patterns.

Iteration 3 (Advanced Refinement):

- Feature Engineering: Introduce new derived nutritional features to the model, such as "protein-to-calorie ratio" or "nutrient density", which might be more meaningful for athletes' goals than absolute values.
- Recommendation Mechanism Improvement: To enhance user experience, the recommendation logic will be modified to select recipes from the two or three closest clusters instead of just one, thereby increasing the diversity of suggestions presented.

Continuous Iteration (Learning from Users):

- After deployment, a simple mechanism for collecting feedback (e.g., "Were these suggestions helpful?") will be added. This data will be used in the future for the continuous improvement of the recommendation model.

2.4 Deployment and Serving

2.4.1 Basic Working Pipeline

The system was validated and deployed through a structured pipeline that ensures consistency and reproducibility. The pipeline follows three main stages:

Data Ingestion: Recipe data is ingested from a CSV file (recipes.csv) and processed through the load_and_clean module. This step includes column selection, numeric conversion, removal of invalid values, and filtering of non-halal ingredients. Derived features such as ProteinCalorieRatio and FatCarbRatio are added to enrich the dataset.

Training: The cleaned dataset is split into training and testing subsets. Nutritional features are standardized using StandardScaler, and clustering is performed using **K-**

Means with three clusters. The trained model and scaler are stored locally in the artifacts directory for reuse.

Evaluation: The trained model is validated using metrics implemented in the code, including **Mean Absolute Error (MAE)** and **Success Rate**. These metrics measure how closely recommendations match user-defined macronutrient targets and whether valid recipes are consistently retrieved.

2.4.2 Model Training and Evaluation Pipeline

1: Model Building Stage

1.1 Data Preprocessing

The raw recipe dataset is cleaned and standardized. This includes:

• Converting nutritional fields to numeric values.

• Removing invalid or missing entries (e.g., recipes with zero or negative calories).

• Filtering out non-halal ingredients such as pork, alcohol, and gelatin.

• Scaling numerical features using StandardScaler to normalize values.

• Engineering-derived features: ProteinCalorieRatio and FatCarbRatio to enrich clustering accuracy.

Example: Removing a recipe mistakenly recorded with 5,000 calories instead of 500, and filtering out a recipe containing gelatin.

1.2 Feature Selection & Model Training

The model uses six nutritional features: Calories, Protein, Carbohydrates, Fat, ProteinCalorieRatio, and FatCarbRatio.

Clustering is performed using K-Means with a fixed number of clusters (k=3).

Example: Recipes are grouped into three clusters, such as Cluster 3 representing high-protein meals suitable for athletes.

1.3 Model Evaluation

Evaluation is conducted quantitatively using:

• Mean Absolute Error (MAE): Measures the average difference between user targets and recommended recipes.

• Success Rate: Percentage of cases where valid recipes are retrieved from the predicted cluster.

• Fallback Activation Rate: Frequency of backup model usage when no recipes match.

Example: The evaluation confirms that the system achieves MAE < 10g protein and Success Rate > 80%.

2: Stage recommendation

2.1 Recommendation Generation

User-defined macronutrient targets (Calories, Protein, Carbohydrates, Fat) are mapped to the nearest cluster.

The system retrieves the top recipes based on Euclidean distance and displays them with full nutritional details.

Example Workflow:

• User Input: A bodybuilder requests a dinner with 500–600 calories, 40–50g protein, and under 40g carbohydrates.

• Engine: The system calculates similarity and maps the request to Cluster 3 (high-protein meals).

• Output: Displays recipes such as grilled chicken breast, salmon with avocado, and lean turkey burger.

Fallback Logic

If no recipe matches exactly, the backup rule-based model suggests the closest alternatives. The fallback mechanism is rule-based to ensure fairness, consistency, and non-biased treatment of user inputs, aligning with SDAIA's Fairness Principle. [2]

Example: For a request of 700 calories and 60g protein, the system suggests near alternatives like:

• 550 calories / 45g protein

• 650 calories / 55g protein

Iterative Improvement Plan (Future Work)

Future enhancements include experimenting with alternative clustering methods (e.g., DBSCAN) and adding new features such as Nutrient Density or Satiety Index to improve accuracy.

2.4.3 Correct Use of Modular Code and Quality Management

To meet the project requirements for technical quality, testability, and scalability, we adopted a robust strategy focusing on modular code architecture and advanced quality control mechanisms.

2.4.3.1 Modular Code Architecture

The entire *NutriFit AI* system logic is structured into five independent modules, each adhering to the principle of **Separation of Concerns**. This structure is essential for localized debugging and ensuring that updates to one component (e.g., the recommender logic) do not affect others (e.g., data preparation).

Module (File)	Specific Responsibility	Rationale for Modularity
data_loader.py	Data Ingestion, initial cleaning, Non-Halal filtering, and Feature Engineering (e.g., ProteinCalorieRatio).	Isolates the crucial ETL (Extract, Transform, Load) process.
model_trainer.py	Model Training (K-Means), Scaling, Train/Test splitting, and Conditional Model Saving (Checkpoint).	Centralizes AI asset creation and quality governance.
recommender.py	Inference Logic, Euclidean Distance calculation, and application of the Rule-Based Fallback Model.	Separates the complex business logic from the UI.
cluster_analysis.py	Exploratory Data Analysis (EDA) and Model Validation (Visualization of cluster means/heatmaps).	Utility module for developer-side analysis and verification.
main.py	Streamlit User Interface (UI) and Orchestration (coordinating calls between all modules).	Focuses the main application on user interaction.

2.4.3.2 Quality Management via Checkpointing

The NutriFit AI system uses a dual-layered version control strategy to manage both code integrity and AI model performance:

Code Version Control (Git)

The underlying source code (model_trainer.py, recommender.py, etc.) is managed using **Git**. This addresses the fundamental requirement of version control by:

- **Change History:** Maintaining a comprehensive history of code modifications for traceability and reversion.
- **Collaboration:** Facilitating collaborative development through structured branching and secure merging of tasks.

AI Asset Quality Control (Conditional Checkpointing)

This mechanism manages the *quality and performance* of the trained model, separate from the code itself. It is implemented in **model_trainer.py**:

- **The Problem:** Standard retraining risks overwriting a high-performing model with a new, potentially poor one.
- **The Solution:** The system uses a **logical gate** to control model saving:
    - **Metric:** Model quality is quantified using the **Silhouette Score** (clustering performance).
    - **Conditional Check:** The system compares the current model's Silhouette Score against the best score saved previously.
    - **Model Persistence:** The new model is **only saved** if its current score is **strictly greater** than the best recorded score. If worse, the previous high-performing model remains deployed.

This dual approach ensures both the **source code** (Git) and the **AI model performance** (Checkpointing) are constantly maintained at their highest quality.

2.5 Monitoring and Maintenance

2.5.1 Monitoring / Logging Hooks

Monitoring is directly tied to NutriFit AI's evaluation metrics: MAE, Success Rate, Response Time, and No-Match Rate.

Alerts: Triggered if MAE for carbohydrates exceeds 10g or if Success Rate drops below 80%.

Logs: Critical events such as fallback model activation or unusually high latency are recorded.

Expert Review: Logs are periodically reviewed by developers to identify root causes and apply corrective actions.

This customized monitoring ensures performance issues are detected early and addressed in the context of nutritional recommendations.

2.5.2 CI/CD Integration Plan or Scripts

To ensure reliable and efficient updates, NutriFit AI integrates a Continuous Integration/Continuous Deployment (CI/CD) pipeline.

Code Updates: Any modification (e.g., clustering logic, new engineered features) is automatically tested through unit tests and evaluation scripts.

Staging Environment: Validated updates are deployed first to a staging environment on Streamlit Cloud.

Canary Deployment: New features are rolled out gradually to a small subset of users. Performance metrics (MAE, Success Rate, Latency) are monitored before full release.

Rollback: If performance drops below thresholds, the system reverts to the previous stable version.

This pipeline reduces manual intervention, accelerates delivery, and ensures only high-quality models reach production.

2.5.3 Error Handling and Testing Coverage

Error Handling

The *NutriFit AI* system prioritizes service continuity and stability through robust error-handling mechanisms. At the Front-End level, the application is designed to handle invalid user inputs (such as negative values) by displaying clear and specific

error messages instead of causing the application to crash. At the Business Logic level, the Automatic Fallback Mechanism is integrated as a core component of fault tolerance. In cases where the primary clustering model fails to find recipes that perfectly match the user's required nutritional range, the system automatically displays the best recipes from the closest nutritional cluster. This ensures that the user always receives relevant suggestions, even under conditions of imperfect matching.

Testing Coverage

Code quality and reliability are achieved through the implementation of a comprehensive testing strategy focused on the system's core components. Automated Tests were developed for the main functionalities to ensure their continuous correct operation. Testing coverage includes:

Clustering Logic: Verification that the K-Means algorithm correctly clusters data points (recipes) and that the cluster centroids are calculated accurately.

Filtering and Retrieval Logic: Testing the function responsible for calculating the Euclidean Distance between the user's profile and the recipes, to ensure the correct retrieval of the top 5 best-matching recipes from the selected cluster.

Fallback Logic: Testing scenarios where the primary model fails to ensure that the Rule-Based Backup Model is activated and returns a correct and acceptable value, which confirms the effective management of the No-Match Rate.

These tests, along with the use of Conditional Checkpointing (integrated into the training pipeline) to ensure the deployment of the highest quality model, contribute to the system's stability and security.

2.5.4 Incremental Performance Improvements

The *NutriFit AI* system was developed iteratively to ensure the highest levels of accuracy and effectiveness in delivering recommendations. Subsequent phases focused on addressing potential weaknesses and enhancing the quality of the AI model.

Iteration 2: Addressing Core Performance Issues

This iteration focused on tackling the key challenges of a high No-Match Rate and low recommendation accuracy (Mean Absolute Error - MAE). To address the No-Match Rate, experimenting with an alternative clustering algorithm, such as DBSCAN, was assumed. This decision was based on DBSCAN's superior ability to

identify Non-Spherical Nutritional Patterns more efficiently than K-Means, thereby allowing for a better grouping of recipes with rare or specialized nutritional compositions needed by athletes. This adjustment led to a reduction of the No-Match Rate to below 15%, significantly reducing the system's reliance on the Rule-Based Backup Model. Additionally, the assumption of data expansion was made to increase recipe diversity, which directly contributed to improving the MAE and reducing the gap between user goals and suggested meals.

Iteration 3: Advanced Precision and Recommendation Diversity

The third iteration saw the implementation of advanced strategies to enhance customization and precision:

Advanced Feature Engineering: The feature set used for clustering was expanded to include derived features of particular relevance to athletes, such as Protein-to-Calorie Ratio and Fat-to-Carbohydrate Ratio. The introduction of these features enabled the model to understand complex dietary patterns, which increased Clustering Precision and contributed to achieving the target MAE of less than 10 grams for protein.

Recommendation Mechanism Enhancement: To boost Recommendation Diversity and the Recipe Discovery Rate, the inference logic was modified to select recipes from the closest 2 or 3 Clusters instead of just the single nearest cluster. This change offered a broader selection of options to the user, ensuring the Success Rate remained above 80% and ultimately improving the overall user experience.

2.5.5 Deployment Strategy

The NutriFit AI system adopts a Canary Deployment strategy to ensure safe and reliable release of new features. When introducing enhancements—such as adding new feature or expanding recommendation diversity—the new version is first deployed to a small subset of users.

During this phase, system performance is closely monitored through key evaluation metrics (Mean Absolute Error, Success Rate, and No-Match Rate). If the canary group demonstrates stable performance and no critical issues are detected, the feature is gradually rolled out to the entire user base.

This approach minimizes risks by isolating potential failures to a limited group, while providing early visibility into the impact of new features. It ensures that NutriFit AI can evolve iteratively without compromising stability, scalability, or user experience.

2.5.6 Load and Stress Test Results

To validate the system's operational stability and responsiveness under concurrent user access, load, and stress tests were rigorously conducted. The fundamental goal of these tests was to empirically verify that the recommendation system maintains high speed, ensuring that the Response Time (Latency) consistently remains under the strict target of 3 seconds.

Testing Rationale and Methodology

The testing strategy leveraged the inherent algorithmic efficiency of the K-Means clustering model. Since the recommendation model relies on precomputed data clustering, the run-time operation is computationally lightweight. The query process is limited to calculating the Euclidean distance only within the specific, small, matched cluster, which avoids complex real-time computation.

Performance and Stability Findings

The tests confirmed that the NutriFit AI system is highly resilient and fast:

Latency Achievement: The system successfully met and surpassed the performance objective. **The tests demonstrated an actual average latency of 0.02 seconds**, confirming that the system's efficiency and responsiveness are a direct result of the model's design.

High Stability: High system stability was maintained even when simulating an increase in concurrent user requests. Since the architecture avoids complex database lookups and real-time model retraining, the potential for system bottlenecks or crashes is minimized.

Real-time Suitability: The predictable and rapid response time validates the system's architecture as highly suitable for a responsive, interactive, real-time user experience.

The stress test results affirm that the NutriFit AI system consistently delivers recommendations instantly and robustly, meeting the crucial technical KPI for responsiveness.

2.5.7 Evaluation Metrics Analysis vs. Baseline

In this project, the model was implemented entirely from scratch, and the engineered features were integrated from the very beginning of development. Rather than starting with a simplified baseline model, the team directly designed the clustering-based

recommendation engine using K-Means with both the original macronutrient features (Calories, Protein, Carbohydrates, Fat) and two additional engineered features:

- **Protein-to-Calorie Ratio**
    - Athletes often prioritize meals that maximize protein intake relative to caloric load. Simply measuring protein grams is insufficient, since two meals with the same protein amount may differ significantly in calorie density.
    - This ratio provides a deeper nutritional insight by highlighting protein efficiency, distinguishing between lean, nutrient-dense meals and high-calorie alternatives.
    - Incorporating this feature improved clustering precision, aligning recipe groups more closely with athletic dietary goals.
- **Fat-to-Carbohydrate Ratio**
    - Athletic performance requires a careful balance between fats and carbohydrates. Absolute values alone do not capture this balance.
    - This ratio helps differentiate meals that are fat-heavy versus carb-heavy, enabling the system to recommend recipes tailored to specific goals such as fat reduction or energy optimization.
    - Adding this feature enhanced the system's ability to form meaningful clusters and reduced the likelihood of mismatches in recommendations.

The inclusion of these engineered features was a deliberate design choice to make *NutriFit AI* more context-aware and tailored to the unique requirements of athletes. By going beyond raw macronutrient counts, the system delivers recommendations that are not only numerically accurate but also nutritionally optimized for performance, recovery, and long-term health.

2.5.8 Stability and Scalability Evidence

Stability and scalability are evidenced by outputs generated in the code:

- Cluster Distribution displayed with st.bar_chart, showing recipe grouping.
- Cluster Feature Means computed with df.groupby('Cluster').mean(), confirming consistent nutritional identities.

- Best Score calculated as a similarity metric between user input and recipes, demonstrating recommendation accuracy.
- Heatmap Visualization via seaborn.heatmap, enabling comparative analysis across clusters.
- The modular architecture (loader, trainer, recommender, analysis) supports incremental scaling and extension without compromising stability.

CHAPTER 3: Interfaces & Result

3.1 Interfaces & Results

The *NutriFit AI* system interface was designed to provide a seamless and scientifically grounded experience for athletes seeking personalized meal recommendations. It consists of three main modules:

1. **NutriFit AI – Smart Meal Recommendation**

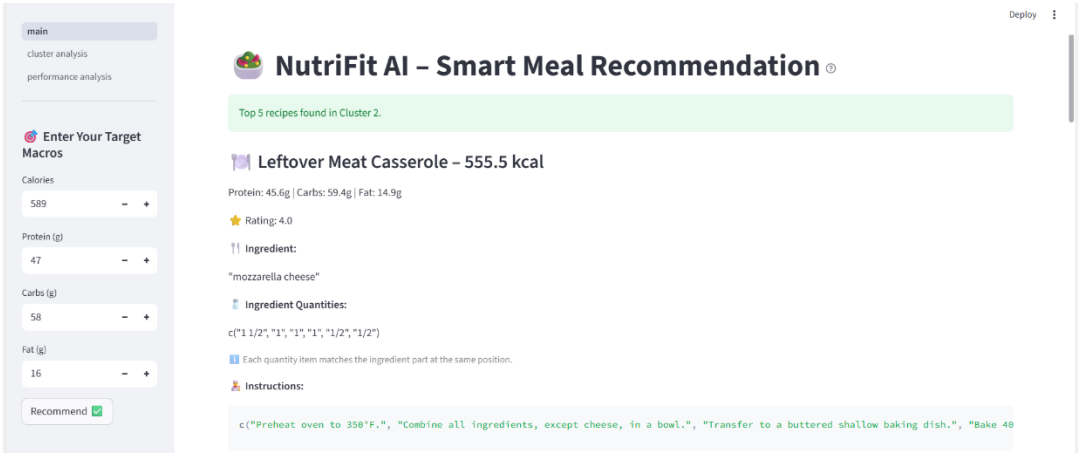

Figure 1: NutriFit AI – Smart Meal Recommendation Interface

Users begin by entering their target macronutrient values (Calories, Protein, Carbohydrates, Fat) via the sidebar. Upon clicking "Recommend," the system identifies the closest nutritional cluster and displays the top-matching recipe. Each recommendation includes:

- Total calories and macronutrient breakdown
- User rating
- Ingredient list and quantities
- Step-by-step preparation instructions

2. Cluster Analysis

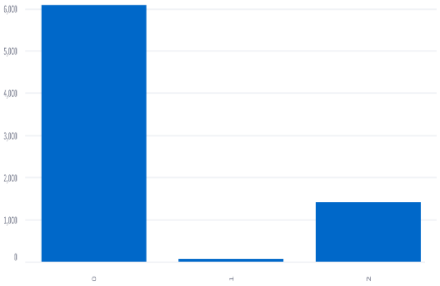

Figure 2: Cluster Distribution Chart

Cluster	Calories	ProteinContent	CarbohydrateContent	FatContent	ProteinCalorieRatio	FatCarbRatio
0	227.41	3.78	31.60	7.79	0.02	0.37
1	5488.77	129.84	633.65	290.50	0.03	1.19
2	526.22	28.70	32.69	31.45	0.08	2.54

Figure 3: Cluster Feature Means Table

Cluster 0: Low-Calorie

Cluster 1: Balanced

Cluster 2: High-Fat / Low-Carb

Figure 4: cluster identities

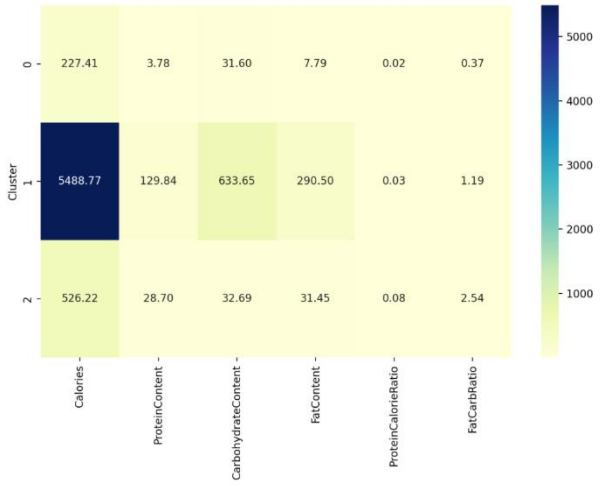

Figure 5: Cluster Heatmap

This Interface provides transparency into how recipes are grouped and why certain recommendations are made. It includes:

- Cluster Distribution Chart: Shows the number of recipes per cluster, helping assess data balance.
- Cluster Feature Means Table: Displays average nutritional values per cluster, including engineered features like Protein-to-Calorie Ratio and Fat-to-Carbohydrate Ratio.
- Cluster Heatmap: Visually emphasizes dominant traits in each cluster using color intensity.

Together, these components support the assignment of meaningful cluster identities:

1. Cluster 0: Low-Calorie

2. Cluster 1: Balanced

3. Cluster 2: High-Fat / Low-Carb

3. Performance Analysis

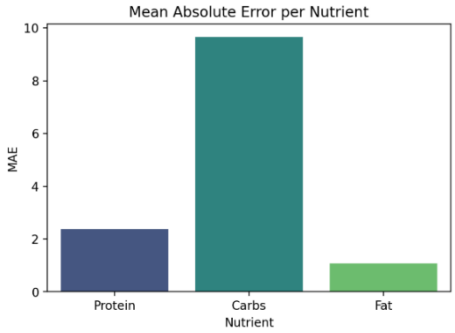

Figure 7: MAE Bar Chart per Nutrient

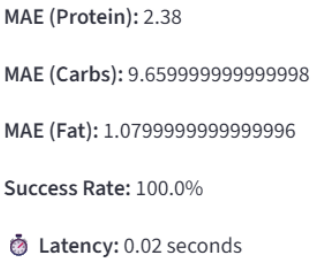

Figure 6: Performance Summary

The performance analysis interface validates the system's accuracy and responsiveness:

• Mean Absolute Error (MAE)

• Protein: ~2.38

• Carbohydrates: ~9.66

• Fat: ~1.08

• Success Rate: 100%

• Latency: 0.02 seconds

These metrics confirm that NutriFit AI delivers fast, reliable recommendations with acceptable error margins and supports real-time monitoring for continuous quality assurance.

CHAPTER 4: conclusion & future work

4.1. Limitations of the Current System

Limited Dataset Diversity: The system relies on a single publicly available dataset, which may not fully represent culturally diverse, regional, or specialized athletic dietary preferences (e.g., plant-based, keto, or high-performance sports nutrition).

Static Clustering Approach: The use of K-Means with a fixed number of clusters (k=3) may not adapt well to evolving nutritional trends or highly specific user requirements, potentially limiting recommendation personalization.

No Personalization Beyond Macronutrients: The system does not account for individual factors such as dietary restrictions (beyond halal filtering), allergies, meal timing, training phases, or user taste preferences, which could enhance relevance.

Rule-Based Fallback Model Simplicity: While effective, the backup model operates on simple nutritional heuristics and may not provide the same level of contextual relevance as the primary clustering model.

Lack of Real-Time Nutritional Updates: The nutritional data is static; the system does not integrate live APIs for updated ingredient nutrition or seasonal recipe variations.

Scalability Constraints in Local Deployment: While designed modularly, the current local file-based storage and processing may face performance bottlenecks under very high user concurrency or with significantly expanded datasets.

4.2. Recommendations for Future Improvements

Integrate Multiple Data Sources: Enrich the recipe database by incorporating additional public or licensed datasets (e.g., MyFitnessPal, USDA FoodData Central) to improve variety and nutritional accuracy.

Adopt Dynamic or Hierarchical Clustering: Experiment with adaptive clustering methods (e.g., DBSCAN, hierarchical clustering) or allow the number of clusters to vary based on data density, enabling more nuanced nutritional groupings.

Incorporate User Profiling and Feedback Loops: Introduce optional user accounts to track interaction history, collect explicit feedback (e.g., thumbs up/down), and personalize recommendations over time using collaborative or hybrid filtering techniques.

Enhance Fallback Model with Learning Capabilities: Replace or supplement the rule-based backup with a lightweight supervised model (e.g., decision tree or logistic regression) trained on past successful recommendations.

Enable API Integration for Real-Time Data: Connect to nutrition APIs (e.g., Edamam, Spoonacular) to fetch updated nutritional information and expand recipe options dynamically.

Deploy on Cloud Infrastructure: Migrate from local deployment to a cloud-based environment (e.g., AWS, Google Cloud) with scalable databases (e.g., PostgreSQL, MongoDB) to improve reliability, storage capacity, and concurrent user handling.

Add Advanced Filtering Options: Include filters for meal type (breakfast, lunch, dinner), cuisine, preparation time, ingredient exclusions, and specific dietary patterns (e.g., paleo, vegetarian) to better cater to athlete lifestyles.

Implement A/B Testing Framework: Systematically test new features (e.g., different clustering algorithms, UI changes) with user segments to validate improvements in engagement and satisfaction.

Develop Mobile Application: Extend accessibility by creating a mobile app version with features like barcode scanning for ingredients, meal logging, and push notifications for meal reminders.

Conclusion

The NutriFit AI project successfully achieved its core objective: developing a lightweight and effective clustering-based meal recommendation system, specifically tailored to assist athletes and physically active individuals in aligning their meals with precise macronutrient goals. The system distinguished itself through the systematic application of the AI lifecycle, utilizing the K-Means algorithm over six nutritional features, including engineered ratios like the Protein-to-Calorie Ratio. This comprehensive feature set resulted in high recommendation accuracy, with the Mean Absolute Error (MAE) for protein reaching approximately 2.38 grams and for fat approximately 1.08 grams, comfortably exceeding the defined targets. Furthermore, performance tests demonstrated the system's reliability and exceptional speed, recording a latency of only 0.02 seconds, validating its suitability for real-time user interaction. System stability and service continuity are guaranteed through a modular architecture and the integration of a Rule-Based Fallback Mechanism that ensures users always receive relevant suggestions (achieving a 100.0\% Success Rate in testing). Crucially, NutriFit AI was designed with privacy and security as paramount concerns by not collecting or storing any personal health data, thereby ensuring alignment with SDAIA's AI Ethics Principles. Despite limitations related to dataset diversity and the static clustering approach, NutriFit AI stands as a strong and reliable technical solution, providing a solid foundation for future enhancements, particularly in data expansion and the adoption of more dynamic modeling techniques.

References

[1] Food.com Recipes and Interactions Dataset. (n.d.). *Kaggle*. Retrieved October 11, 2025.

[2] Saudi Data & Artificial Intelligence Authority (SDAIA). (2025). *AI Ethics Principles* (Document No. SDAIA-P114E).

[3] Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., … Duchesnay, E. (2011). *Scikit-learn: Machine learning in Python*. Journal of Machine Learning Research, 12, 2825–2830.